

Lecture 6: file I/O, finish Unix

Announcements

- Projects
 - 4A assigned: due in class on Weds
 - 2B prompt available Saturday AM, due on Friday
- No class on Weds April 27
 - likely a short YouTube replacement lecture

Outline

- Review
- File I/O
- Project 2B
- Redirection
- Pipes

Outline

- Review
- File I/O
- Project 2B
- Redirection
- Pipes

Makefile example: multiplier lib



```
C02LN00GFD58:code hank$ cat Makefile
lib: doubler.o tripler.o
      ar r libmultiplier.a doubler.o tripler.o
      cp libmultiplier.a ~/multiplier/lib
      cp multiplier.h ~/multiplier/include

doubler.o: doubler.c
          gcc -c doubler.c

tripler.o: tripler.c
          gcc -c tripler.c
C02LN00GFD58:code hank$
C02LN00GFD58:code hank$
C02LN00GFD58:code hank$ make
ar r libmultiplier.a doubler.o tripler.o
cp libmultiplier.a ~/multiplier/lib
cp multiplier.h ~/multiplier/include
C02LN00GFD58:code hank$
C02LN00GFD58:code hank$ touch doubler.c
C02LN00GFD58:code hank$ make
gcc -c doubler.c
ar r libmultiplier.a doubler.o tripler.o
cp libmultiplier.a ~/multiplier/lib
cp multiplier.h ~/multiplier/include
C02LN00GFD58:code hank$
```

Unix command: tar

- `tar cvf 330.tar file1 file2 file3`
 - puts 3 files (file1, file2, file3) into a new file called 330.tar
- `scp 330.tar @ix:~`
- `ssh ix`
- `tar xvf 330.tar`
- `ls`
`file1 file2 file`

Memory Errors

- Array bounds read

```
int main()
{
    int var;
    int arr[3] = { 0, 1, 2 };
    var=arr[3];
}
```

- Array bounds write

```
int main()
{
    int var = 2;
    int arr[3];
    arr[3]=var;
}
```

Memory Errors

- Free memory read / free memory write

```
int main()
{
    int *var = malloc(sizeof(int)*2);
    var[0] = 0;
    var[1] = 2;
    free(var);
    var[0] = var[1];
}
```

Vocabulary: “dangling pointer”: pointer that points to memory that has already been freed.

Memory Errors

- Freeing unallocated memory

```
int main()
{
    int *var = malloc(sizeof(int)*2);
    var[0] = 0;
    var[1] = 2;
    free(var);
    free(var);
}
```

Memory Errors

- Freeing non-heap memory

```
int main()
{
    int var[2]
    var[0] = 0;
    var[1] = 2;
    free(var);
}
```

Memory Errors

- NULL pointer read / write

```
int main()
{
    char *str = NULL;
    printf(str);
    str[0] = 'H';
}
```

- NULL is never a valid location to read from or write to, and accessing them results in a “segmentation fault”
 - remember those memory segments?

Memory Errors

- Uninitialized memory read

```
int main()
{
    int *arr = malloc(sizeof(int)*10);
    int v2=arr[3];
}
```

Memory error in action

```
fawcett:error child$ cat t.c
#include <stdio.h>

int main()
{
    int *X = NULL;
    printf("X is %d\n", *X);
}
fawcett:error child$ gcc t.c
fawcett:error child$ ./a.out
Segmentation fault
fawcett:error child$ █
```

Project 4A

- Posted now
- You will practice debugging & using a debugger
 - There are 3 programs you need to debug
 - In this case, “debug” means identify the bug
 - Does not mean fix the bug
 - Can use gdb or lldb
 - May want to run on ix
- Worksheet due in class next week

ASCII Character Set

ASCII Code Chart

0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
0 NUL	SOH	STX	ETX	EOT	ENQ	ACK	BEL	BS	HT	LF	VT	FF	CR	SO	SI
1 DLE	DC1	DC2	DC3	DC4	NAK	SYN	ETB	CAN	EM	SUB	ESC	FS	GS	RS	US
2 !	"	#	\$	%	&	'	()	*	+	.	-	.	/	
3 0	1	2	3	4	5	6	7	8	9	:	;	<	=	>	?
4 @	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
5 P	Q	R	S	T	U	V	W	X	Y	Z	[\]	^	-
6 `	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o
7 p	q	r	s	t	u	v	w	x	y	z	{	}	-	DEL	

There have been various extensions to ASCII ...
now more than 128 characters

Many special characters are handled outside this convention

signed vs unsigned chars

- signed char (“char”):
 - valid values: -128 to 127
 - size: 1 byte
 - used to represent characters with ASCII
 - values -128 to -1 are not valid
- unsigned char:
 - valid values: 0 to 255
 - size: 1 byte
 - used to represent data

character strings

- A character “string” is:
 - an array of type “char”
 - that is terminated by the NULL character
- Example:

```
char str[12] = "hello world";
```

 - str[11] = '\0' (the compiler did this automatically)
- The C library has multiple functions for handling strings

Character strings example

```
128-223-223-72-wireless:330 hank$ cat string.c
#include <stdio.h>

int main()
{
    char str[12] = "hello world";
    char *str2 = str+6;

    printf("str is \"%s\" and str2 is \"%s\"\n",
           str, str2);

    str[5] = '\0';

    printf("Now str is \"%s\" and str2 is \"%s\"\n",
           str, str2);
}

128-223-223-72-wireless:330 hank$ gcc string.c
128-223-223-72-wireless:330 hank$ ./a.out
str is "hello world" and str2 is "world"
Now str is "hello" and str2 is "world"
```

memcpy

MEMCPY(3)

BSD Library Functions Manual

MEMCPY(3)

NAME**memcpy** -- copy memory area**LIBRARY**

Standard C Library (libc, -lc)

SYNOPSIS

```
#include <string.h>

void *
memcpy(void *restrict dst, const void *restrict src, size_t n);
```

DESCRIPTION

The **memcpy()** function copies *n* bytes from memory area *src* to memory area *dst*. If *dst* and *src* overlap, behavior is undefined. Applications in which *dst* and *src* might overlap should use **memmove(3)** instead.

RETURN VALUES

The **memcpy()** function returns the original value of *dst*.

I mostly use C++, and I still use memcpy all the time

if-then-else

```
int val = (X < 2 ? X : 2);
```

↔

```
if (X<2)
{
    val = X;
}
else
{
    val = 2;
}
```

Outline

- Review
- File I/O
- Project 2B
- Redirection
- Pipes

File I/O: streams and file descriptors

- Two ways to access files:
 - File descriptors:
 - Lower level interface to files and devices
 - Provides controls to specific devices
 - Type: small integers (typically 20 total)
 - Streams:
 - Higher level interface to files and devices
 - Provides uniform interface; easy to deal with, but less powerful
 - Type: FILE *

Streams are more portable, and more accessible to beginning programmers. (I teach streams here.)

File I/O

- Process for reading or writing
 - Open a file
 - Tells Unix you intend to do file I/O
 - Function returns a “FILE *”
 - Used to identify the file from this point forward
 - Checks to see if permissions are valid
 - Read from the file / write to the file
 - Close the file

Opening a file

- FILE *handle = fopen(filename, mode);

The argument `mode` points to a string beginning with one of the following sequences (Additional characters may follow these sequences.):

``r'' Open text file for reading. The stream is positioned at the beginning of the file.

``r+'' Open for reading and writing. The stream is positioned at the beginning of the file.

Example: FILE *h = fopen("/tmp/330", "wb");

exist, otherwise it is truncated. The stream is positioned at the beginning of the file.

Close when you are done with “fclose”

``a+'' Open for reading and writing. The file is created if it does not exist. The stream is positioned at the end of the file. Subsequent writes to the file will always end up at the then current end of file, irrespective of any intervening fseek(3) or similar.

Note: #include <stdio.h>

acter or
rings
99:1990

Reading / Writing

FREAD(3) BSD Library Functions Manual FREAD(3)

NAME
fread, fwrite -- binary stream input/output

LIBRARY
Standard C Library (libc, -lc)

SYNOPSIS

```
#include <stdio.h>

size_t
fread(void *restrict ptr, size_t size, size_t nitems, FILE *restrict stream);

size_t
fwrite(const void *restrict ptr, size_t size, size_t nitems,
FILE *restrict stream);
```

DESCRIPTION
The function **fread()** reads nitems objects, each size bytes long, from the stream pointed to by stream, storing them at the location given by ptr.

The function **fwrite()** writes nitems objects, each size bytes long, to the stream pointed to by stream, obtaining them from the location given by ptr.

RETURN VALUES
The functions **fread()** and **fwrite()** advance the file position indicator for the stream by the number of bytes read or written. They return the number of objects read or written. If an error occurs, or the end-of-file is reached, the return value is a short object count (or zero).

Example

```
C02LN00GFD58:330 hank$ cat rw.c
#include <stdio.h>
#include <string.h>

int main(int argc, char *argv[])
{
    char *hello = "hello world: file edition\n";
    FILE *f = fopen("330", "w");
    fwrite(hello, sizeof(char), strlen(hello), f);
    fclose(f);
}
C02LN00GFD58:330 hank$ gcc rw.c
C02LN00GFD58:330 hank$ ./a.out
C02LN00GFD58:330 hank$ cat 330
hello world: file edition
```

fseek

```
int  
fseek(FILE *stream, long offset, int whence);
```

The **fseek()** function sets the file position indicator for the stream pointed to by stream. The new position, measured in bytes, is obtained by adding offset bytes to the position specified by whence. If whence is set to SEEK_SET, SEEK_CUR, or SEEK_END, the offset is relative to the start of the file, the current position indicator, or end-of-file, respectively. A successful call to the **fseek()** function clears the end-of-file indicator for the stream and undoes any effects of the ungetc(3) and ungetwc(3) functions on the same stream.

f~~tell~~

```
long  
ftell(FILE *stream);
```

The **f~~tell~~()** function obtains the current value of the file position indicator for the stream pointed to by stream.

We have everything we need to make a copy command...

- fopen
- fread
- fwrite
- fseek
- ftell

Can we do this together as a class?

argc & argv

- two arguments to every C program
- argc: how many command line arguments
- argv: an array containing each of the arguments
- “./a.out hank childs”
- → argc == 3
- → argv[0] = “a.out”, argv[1] = “hank”,
argv[2] = “childs”

O

```
#include <stdio.h>
#include <printf.h>
#include <stdlib.h>

int main(int argc, char *argv[])
{
    FILE *f_in, *f_out;
    int buff_size;
    char *buffer;

    if (argc != 3)
    {
        printf("Usage: %s <file1> <file2>\n", argv[0]);
        exit(EXIT_FAILURE);
    }

    f_in = fopen(argv[1], "r");
    fseek(f_in, 0, SEEK_END);
    buff_size = ftell(f_in);
    fseek(f_in, 0, SEEK_SET);

    buffer = malloc(buff_size);
    fread(buffer, sizeof(char), buff_size, f_in);

    printf("Copying %d bytes from %s to %s\n", buff_size, argv[1], argv[2]);

    f_out = fopen(argv[2], "w");
    fwrite(buffer, sizeof(char), buff_size, f_out);

    fclose(f_in);
    fclose(f_out);

    return 0;
}
```

Return values in shells

```
C02LN00GFD58:330 hank$ ./a.out copy.c copy2.c
Copying 697 bytes from copy.c to copy2.c
C02LN00GFD58:330 hank$ echo $?
0
C02LN00GFD58:330 hank$ ./a.out copy.c
Usage: ./a.out <file1> <file2>
C02LN00GFD58:330 hank$ echo $?
1
```

\$? is the return value of the last executed command

Printing to terminal and reading from terminal

- In Unix, printing to terminal and reading from terminal is done with file I/O
- Keyboard and screen are files in the file system!
 - (at least they were ...)

Standard Streams

- Wikipedia: “preconnected input and output channels between a computer program and its environment (typically a text terminal) when it begins execution”
- Three standard streams:
 - stdin (standard input)
 - stdout (standard output)
 - stderr (standard error)

What mechanisms in C allow you to access standard streams?

printf

- Print to stdout
 - `printf("hello world\n");`
 - `printf("Integers are like this %d\n", 6);`
 - `printf("Two floats: %f, %f", 3.5, 7.0);`

fprintf

- Just like printf, but to streams
- `fprintf(stdout, "helloworld\n");`
 - → same as printf
- `fprintf(stderr, "helloworld\n");`
 - prints to “standard error”
- `fprintf(f_out, "helloworld\n");`
 - prints to the file pointed to by FILE *f_out.

buffering and printf

- Important: printf is buffered
- So:
 - printf puts string in buffer
 - other things happen
 - buffer is eventually printed
- But what about a crash?
 - printf puts string in buffer
 - other things happen ... including a crash
 - buffer is never printed!

Solutions: (1) fflush, (2) fprintf(stderr) always flushed

Outline

- Review
- File I/O
- Project 2B
- Redirection
- Pipes

Project 2B

(which means submitted by 6am on April 24th, 2015)

Worth 4% of your grade

Assignment: Write a program that reads the file "2E_binary_file". This file contains a two-dimensional array of integers, that is 10x10. You are to read in the 5x5 bottom left corner of the array. That is, the values 0-4, 10-14, 20-24, 30-34, and 40-44. You may only read 25 integers total. Do not read all 100 and throw some out. You will then write out the new 5x5 array. Please write this as strings, one integer per line (25 lines total). You should be able to "cat" the file afterwards and see the values.

Use Unix file streams for this project (i.e., fopen, fread, fseek, fprintf). Your program will be checked for good programming practices. (Close your file streams, use memory correctly, etc. I am not referring to style, variable initialization, etc.)

Also, add support for command line arguments (argc and argv).

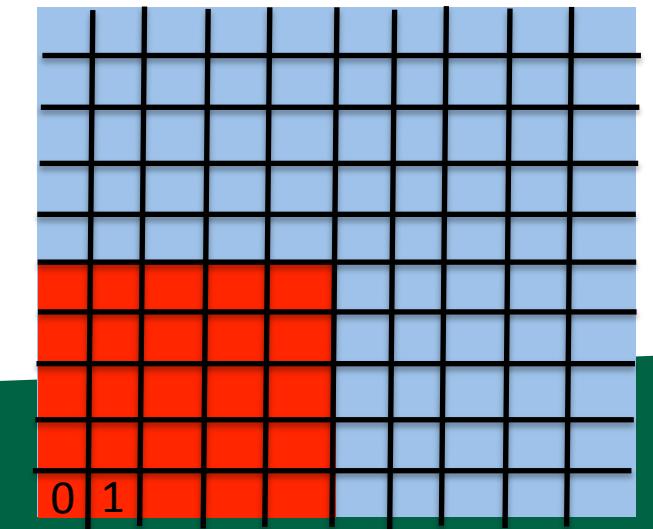
Your program should run as:

`./<prog_name> <input_name> <output_name>`

(The input_name will be 2E_binary_file, unless you change it.)

Finally, note that I am handing you a binary file. I think we are all little endian, and so it will be fine. But, if it is big endian, then we will have a problem. You can check if it is little endian by printing the first two values of the file. They should be "0" and "1".

Please submit a tarball with (1) a Makefile (should be simple), (2) your source code, and (3) the output ASCII file from running your program, with the name "ASCII_output".



Outline

- Review
- File I/O
- Project 2B
- **Redirection**
- Pipes

Unix shells allows you to manipulate standard streams.

- “>” redirect output of program to a file
- Example:
 - ls > output
 - echo “this is a file” > output2
 - cat file1 file2 > file3

Unix shells allows you to manipulate standard streams.

- “<” redirect file to input of program
- Example:
 - `python < myscript.py`
 - Note: python quits when it reads a special character called EOF (End of File)
 - You can type this character by typing Ctrl-D
 - This is why Python quits when you type Ctrl-D
 - (many other programs too)

Unix shells allows you to manipulate standard streams.

- “>>” concatenate output of program to end of existing file
 - (or create file if it doesn’t exist)
- Example:
 - echo “I am starting the file” > file1
 - echo “I am adding to the file” >> file1
 - cat file1

I am starting the file

I am adding to the file

What's happening here?

```
C02LN00GFD58:330 hank$ mkdir tmp
C02LN00GFD58:330 hank$ cd tmp
C02LN00GFD58:tmp hank$ touch f1
C02LN00GFD58:tmp hank$ ls f1 f2 > out
ls: f2: No such file or directory
C02LN00GFD58:tmp hank$ cat out
f1
```

ls is outputting its error messages to stderr

Redirecting stderr in a shell

```
C02LN00GFD58:Documents hank$ cd ~/330
C02LN00GFD58:330 hank$ mkdir tmp
C02LN00GFD58:330 hank$ cd tmp
C02LN00GFD58:tmp hank$ touch f1
C02LN00GFD58:tmp hank$ ls f1 f2 > out
ls: f2: No such file or directory
C02LN00GFD58:tmp hank$ cat out
f1
C02LN00GFD58:tmp hank$ ls f1 f2 > out 2>out_error
C02LN00GFD58:tmp hank$ cat out_error
ls: f2: No such file or directory
```

Redirecting stderr to stdout

```
C02LN00GFD58:330 hank$ mkdir tmp
C02LN00GFD58:330 hank$ cd tmp
C02LN00GFD58:tmp hank$ touch f1
C02LN00GFD58:tmp hank$ ls f1 f2 > out
ls: f2: No such file or directory
C02LN00GFD58:tmp hank$ cat out
f1
C02LN00GFD58:tmp hank$ ls f1 f2 > out 2>out_error
C02LN00GFD58:tmp hank$ cat out_error
ls: f2: No such file or directory
C02LN00GFD58:tmp hank$ ls f1 f2 > out 2>&1
C02LN00GFD58:tmp hank$ cat out
ls: f2: No such file or directory
f1
```

Convenient when you want both to go to the same stream

Outline

- Review
- File I/O
- Project 2B
- Redirection
- Pipes

c functions: fork and pipe

- fork: duplicates current program into a separate instance
 - Two running programs!
 - Only differentiated by return value of fork (which is original and which is new)
- pipe: mechanism for connecting file descriptors between two forked programs

Through fork and pipe, you can connect two running programs. One writes to a file descriptor, and the other reads the output from its file descriptor.

Only used on special occasions.
(And one of those occasions is with the shell.)

pipes in Unix shells

```
C02LN00GFD58:tmp hank$ cat printer.c
#include <stdio.h>
int main() { printf("Hello world\n"); }
C02LN00GFD58:tmp hank$ cat doubler.c
#include <stdio.h>
int main()
{
    int ch = getc(stdin);
    while (ch != EOF)
    {
        printf("%c%c", ch, ch);
        ch = getc(stdin);
    }
}
C02LN00GFD58:tmp hank$ gcc -o printer printer.c
C02LN00GFD58:tmp hank$ gcc -o doubler doubler.c
C02LN00GFD58:tmp hank$ ./printer | ./doubler
HHeelllloo  wwoorrlldd

C02LN00GFD58:tmp hank$
```

- represented with “|”
- output of one program becomes input to another program

Very useful programs

- grep: keep lines that match pattern, discard lines that don't match pattern

```
C02LN00GFD58:Documents hank$ ls -l | grep ppt
-rw-r--r--@ 1 hank staff 3278589 Apr  5 11:40 CIS330_Lec2.pptx
-rw-r--r--@ 1 hank staff 2220104 Apr  8 20:57 CIS330_Lec3.pptx
-rw-r--r-- 1 hank staff 3899863 Jan 21 09:26 CIS610_lec2.pptx
-rw-r--r-- 1 hank staff 4629257 Jan 30 10:24 CIS610_lec3.pptx
-rw-r--r-- 1 hank staff 21382185 Mar 25 12:40 CIS_colloquium2013.pptx
-rw-r--r-- 1 hank staff 21382185 Jan  7 12:21 CIS_colloquium_2013.pptx
-rw-r--r--@ 1 hank staff 2172179 Dec 20 15:24 ICS_results.pptx
-rw-r--r--@ 1 hank staff 4841050 Nov 13 10:10 MBTI.pptx
-rw-r--r--@ 1 hank staff 2031749 Apr  5 16:20 SC14_flow.pptx
-rw-r--r-- 1 hank staff 17972476 Mar 25 12:43 VMV_2013.pptx
-rw-r--r--@ 1 hank staff 98149068 Apr  1 10:25 aachen.pptx
-rw-r--r-- 1 hank staff 9815146 Feb 24 07:00 childs_poster_SDAV_AHM_2014.pptx
-rw-r--r--@ 1 hank staff 592243 Feb 26 04:09 childs_sdav_slides.pptx
-rw-r--r--@ 1 hank staff 15765504 Feb 13 14:57 cig_exascale.ppt
-rw-r--r--@ 1 hank staff 16699392 Jan  7 12:14 cis610_Lec1.ppt
-rw-r--r-- 1 hank staff 3159872 Jan  7 11:15 epgv_cgf.pptx
-rw-r--r--@ 1 hank staff 15767552 Mar 23 02:48 eu_regional_school.ppt
-rw-r--r--@ 1 hank staff 35099136 Mar 25 09:42 eu_regional_school_part1.ppt
-rw-r--r--@ 1 hank staff 10775552 Mar 25 04:49 eu_regional_school_part1B.ppt
-rw-r--r--@ 1 hank staff 72966144 Mar 26 08:43 eu_regional_school_part2.ppt
-rw-r--r-- 1 hank staff 7571317 Mar 25 12:53 ilm_booth_talk.pptx
```

Very useful programs

- sed: replace pattern 1 with pattern 2
 - sed s/pattern1/pattern2/g
 - s means substitute
 - g means “global” ... every instance on the line

sed is also available in “vi”

:%s/pattern1/pattern2/g (% means all lines)
:103,133s/p1/p2/g (lines 103-133)

Bonus topic: wildcards

- '*' is a wildcard with unix shells

```
fawcett:tmp child$ ls
Abe          Chavarria    Hebb        Macy        Smith
Alajaji     Chen          Jia         Maguire    Steelhammer
Alamoudi    Clark         Kine        Michlanski Szczepanski
Anastas     Collier       Lee         Moreno     Totten
Andrade     Costello      Legge      Olson      Vega-Fujioka
Ballarche   Donnelly     Li          Owen       Wang
Brennan     Etzel        Lin         Pogrebinsky Whiteley
Brockway    Friedrich    Liu         Qin        Woodruff
Brogan      Garvin       Lopes      Rhodes     Xu
Brooks      Gonzales     Luo         Roberts    Yaconelli
Bruce       Guo          Lynch      Rodriguez Young
Carlton     Hampton      Lyon       Roush      Zhang
Chalmers   Harris       Machado   Rozenboim de
fawcett:tmp child$ ls C*
Carlton     Chavarria    Clark      Costello
Chalmers   Chen          Collier
fawcett:tmp child$ ls *z
Rodriguez
fawcett:tmp child$ ls *ee*
Lee          Steelhammer
fawcett:tmp child$ ls *e*ee*
Lee          Legge        Steelhammer Whiteley
```

'?' is a wildcard that matches exactly one character

Other useful shell things

- ‘tab’: auto-complete
- esc=: show options for auto-complete
- Ctrl-A: go to beginning of line
- Ctrl-E: go to end of line
- Ctrl-R: search through history for command